

Refuse to Crash with Re-FUSE

Swaminathan Sundararaman, Laxman Visampalli,
Andrea C. Arpaci-Dusseau, Remzi H. Arpaci-Dusseau




A D S L



Kernel-level File Systems

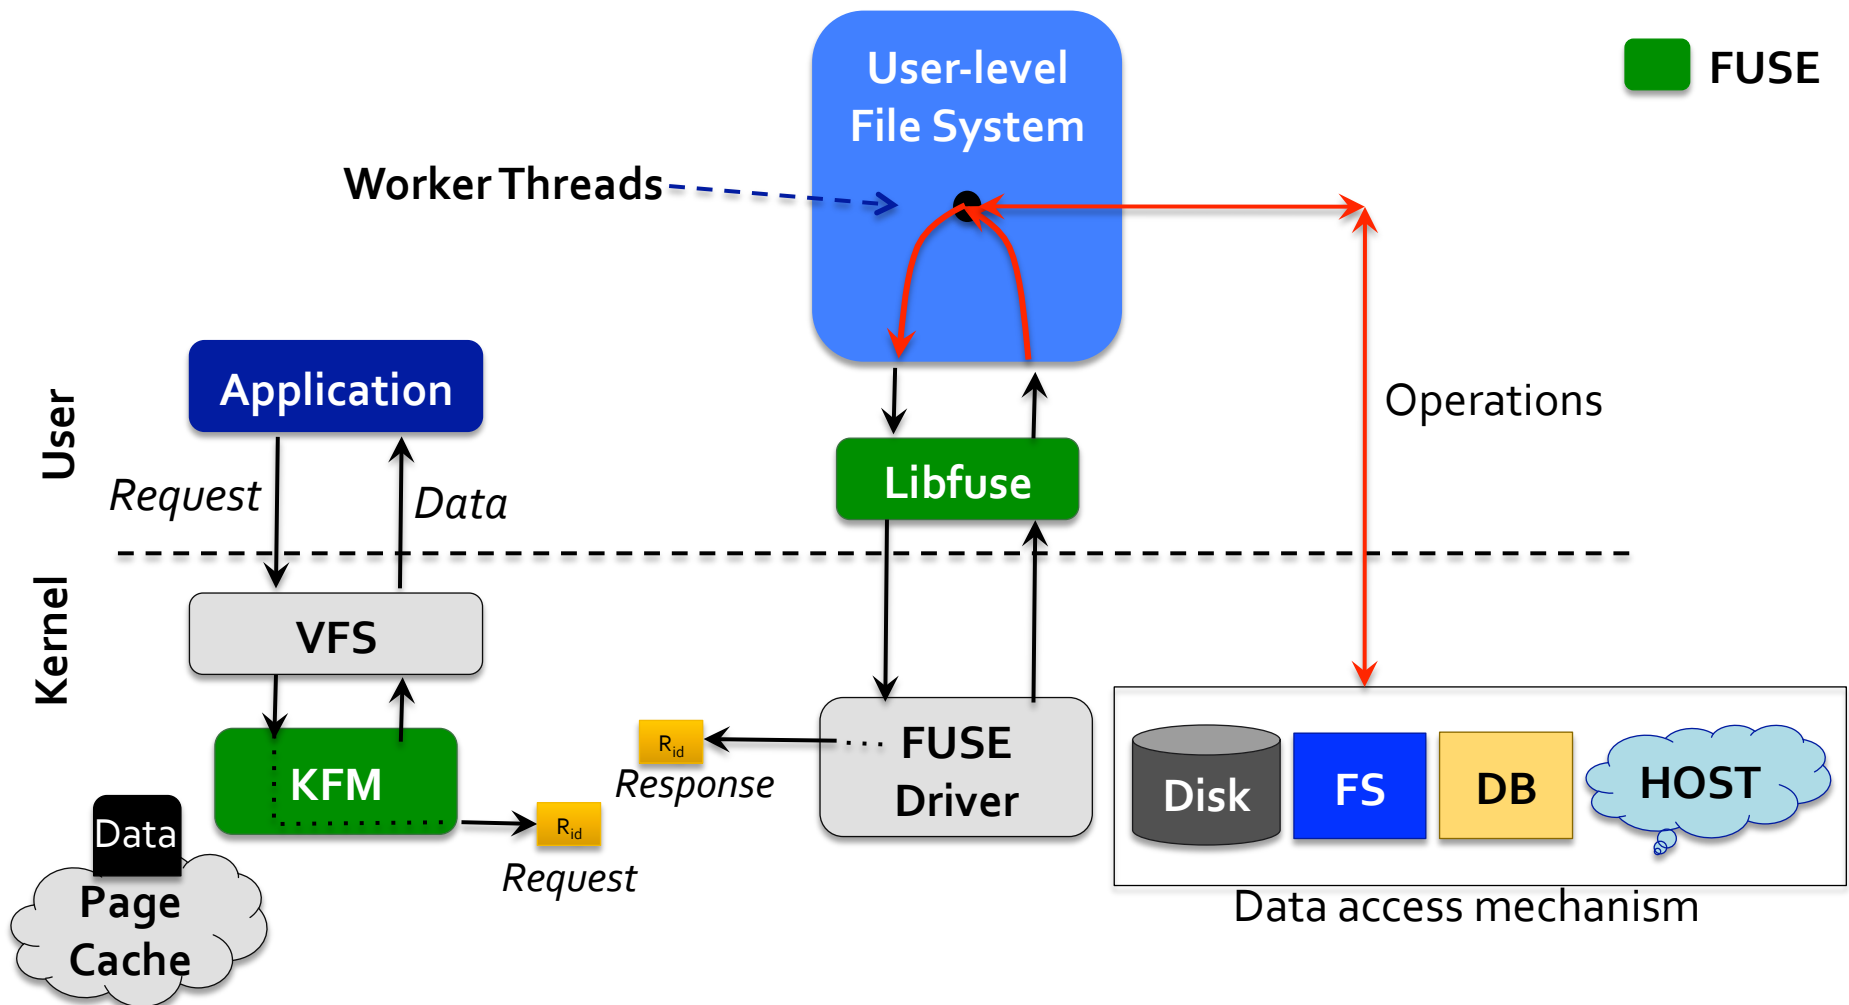
- Hard to develop FSes
 - Interact with OS components
 - Support variety of features
 - Difficult to debug
- *Require:* skilled developers
- New file system in 5-10 years
 - Adding user-desired features is difficult



FUSE: File system in USErspace

- Framework to run FS as a user-level process
 - Simplify development and deployment of FSes
 - FS interface to access underlying data
 - Database, email, ftp, http, ssh
- ~180 user-level file systems in few years
 - Can be customized to do just one thing
 - Compression, remote access (interface to cloud storage)
 - Raw device, pass through, and network-based

FUSE Architecture



User-level File System Issues

- Quick development
 - Not your typical file-system developer
 - No rigorous testing
 - No good documentation on FUSE
 - API, error scenarios
- Result in crashes
 - Require manual intervention to fix the problem
 - Repair & restart
 - Users feels that file system "*does not work*"
 - Decreases adoption chances

Re-FUSE: Restartable FUSE

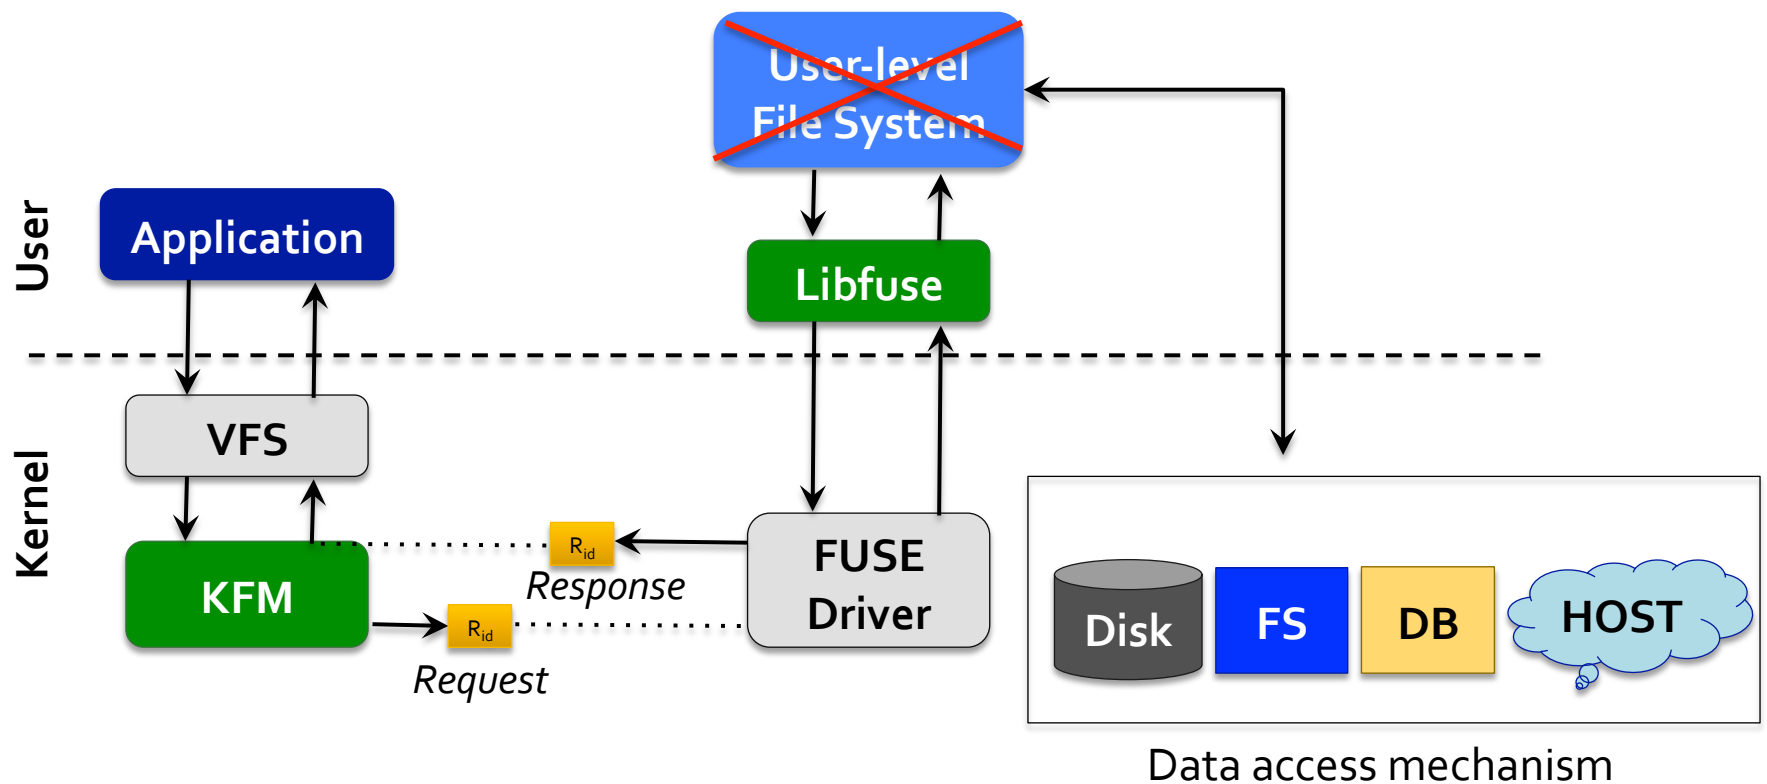
- Framework to restart FS on crashes
 - A generic mechanism to restart in **transparent** and **stateful**
 - Applications are *oblivious* to FS crashes
- Novel techniques
 - Request tagging, system call logging, non-interruptible system calls
 - Performance: page versioning and socket buffer caching
- Evaluation (Linux 2.6.18, FUSE 2.7.4, NTFS-3g, AVFS, and SSHFS)
 - Generality: **< 10 lines of code** for each FS
 - Robustness: **60 controlled & 300 random** fault injection
 - Performance: **< 13%** for both micro & macro benchmarks
 - Recovery time: **< 300 milliseconds** to restart FS

Outline

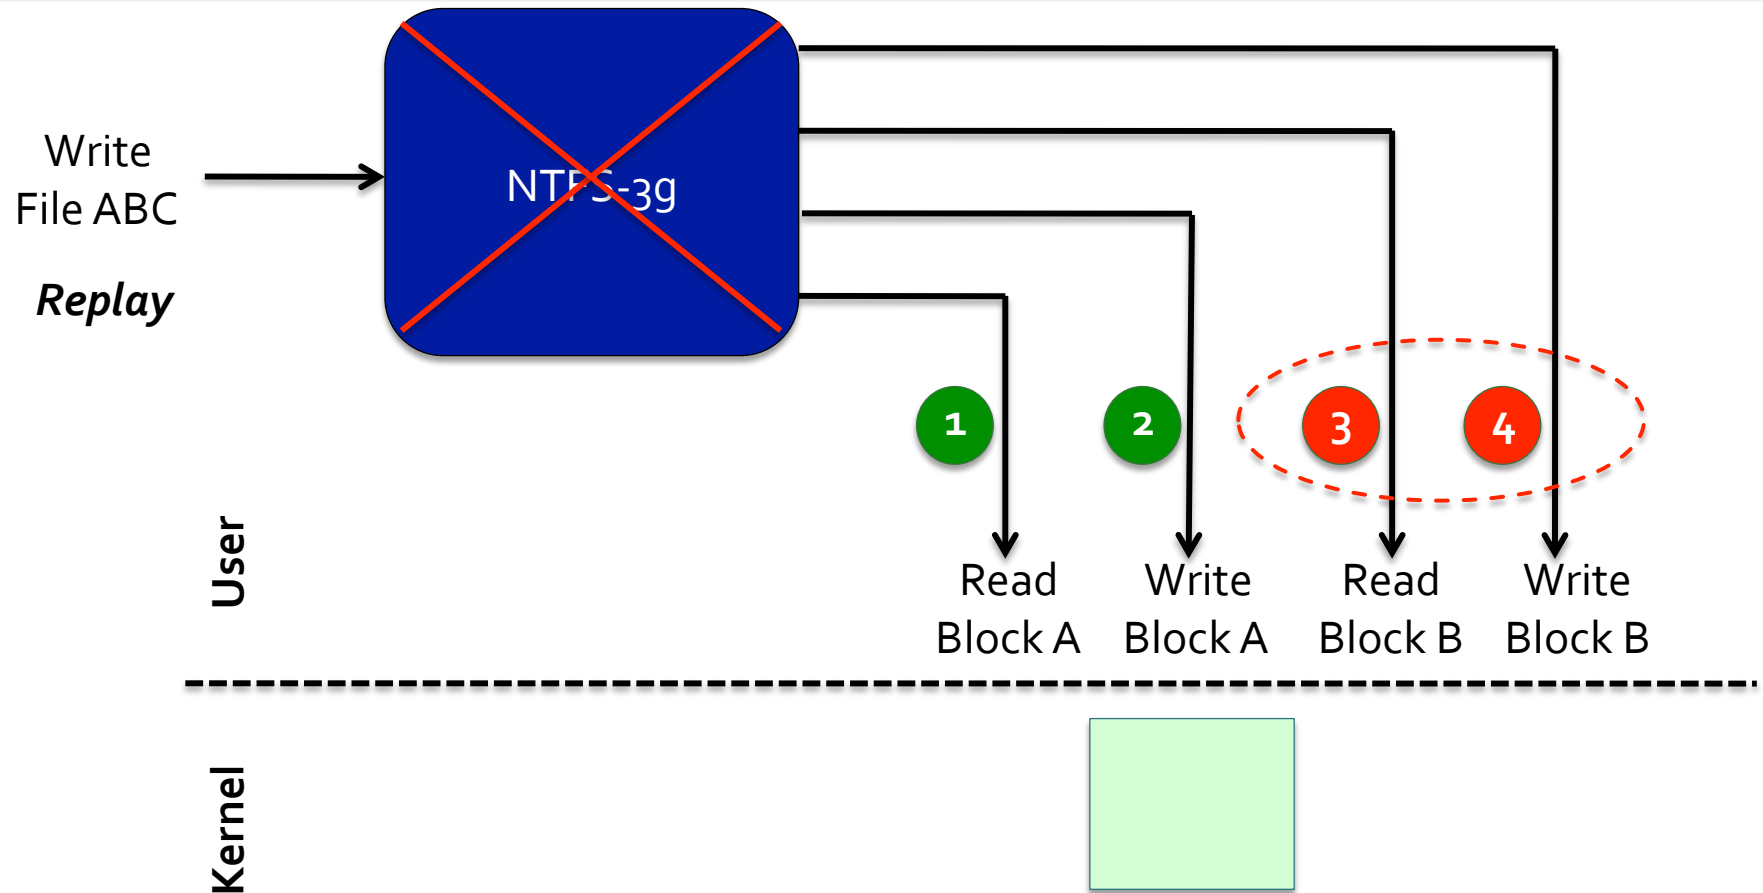
- Background
- Challenges
- Re-FUSE
- Evaluation
- Conclusions

Failure Model

- Only the user-level file system is unavailable

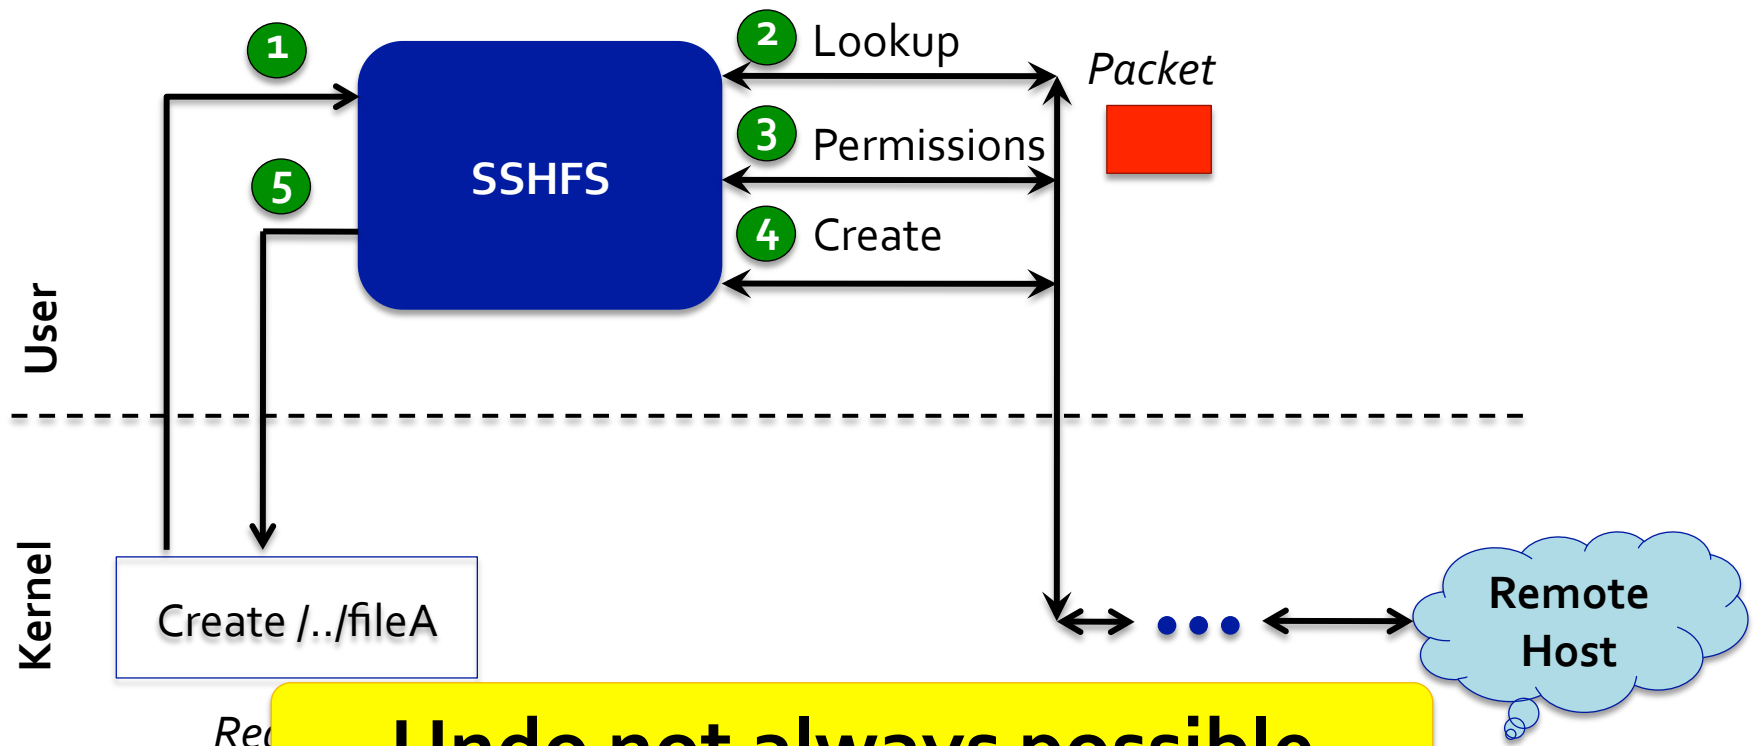


Can Simple Re-execution Work?



File System is left in an inconsistent state

How About Undoing Operations?



Undo not always possible


Need: alternate mechanism to restore FSES

Outline

- Background
- Challenges
- Re-FUSE
- Evaluation
- Conclusions

Re-FUSE


- Restartable FUSE
 - Designed to deal with file system crashes
 - Framework built inside FUSE and OS
- **Principle:** inconsistent \longrightarrow consistent state



Continue from the last completed operation

Components of Re-FUSE

- Fault Detection
 - Fail-stop, transient faults
 - Monitors FS crashes
- Fault Anticipation
 - Records file-system state
- Fault Recovery
 - Restart FS process and re-execute requests




Fault Anticipation

Additional work done in preparation of failure


- Simplified due to FUSE architecture
 - Separate process executes FS requests
 - In-flight requests are preserved
 - Only the user-level file system crashes on failure
 - Applications state and updates (data) are preserved
- **Need:** generic mechanism to track progress
 - Mimic re-execution with sufficient recorded state

Challenges In Tracking Progress

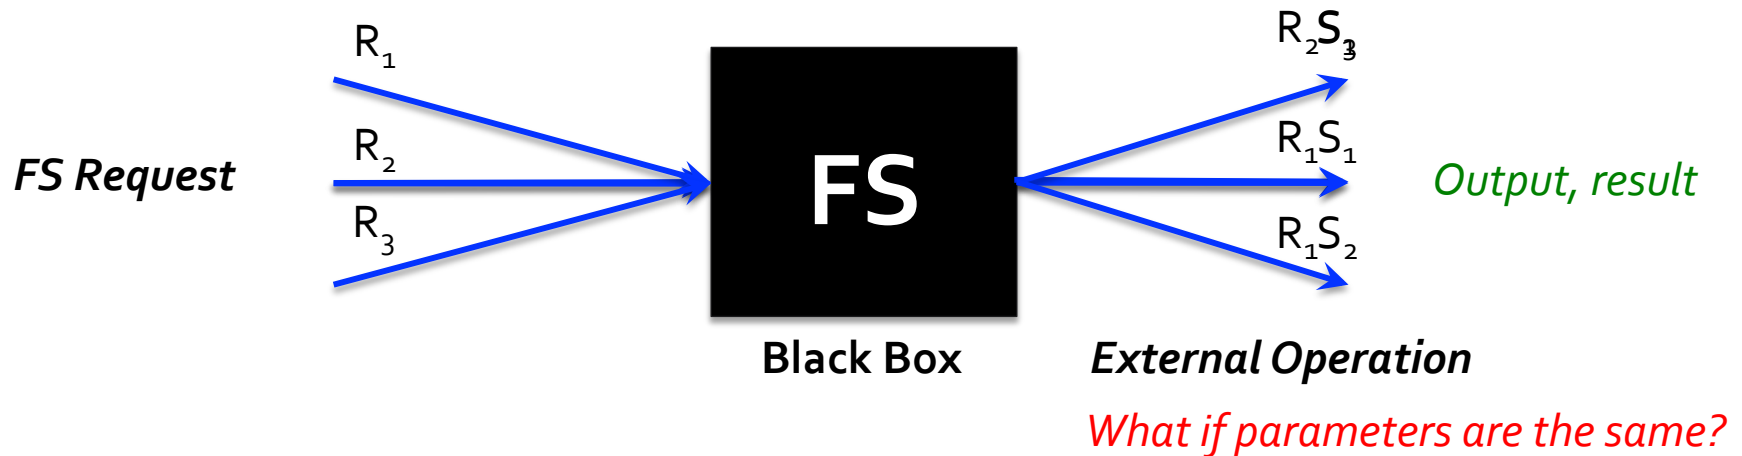
- Correlate FS actions with in-flight requests
 - Decoupled execution & multi threading



- Determine the state of individual sub tasks
 - Request splitting



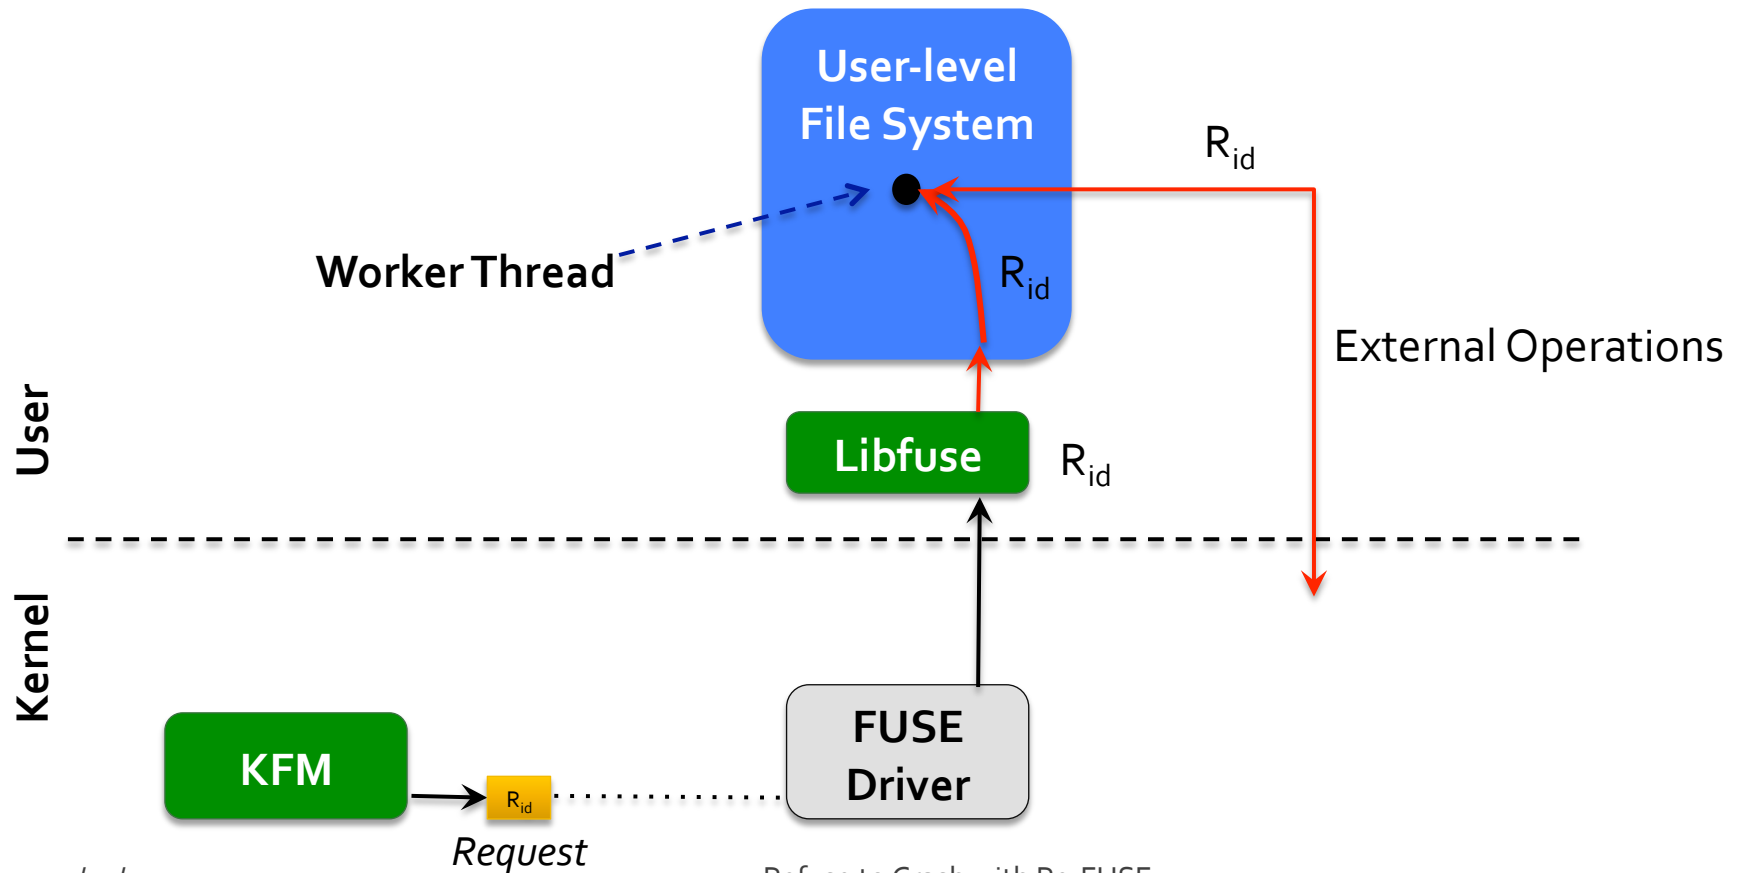
Straw Man Approach



- Single request, single external operation
 - Require no additional support
- Single request, many external operations
 - Sequence number to correlate operations
- Many requests, many external operations
 - Request id and sequence number to correlate request and operation

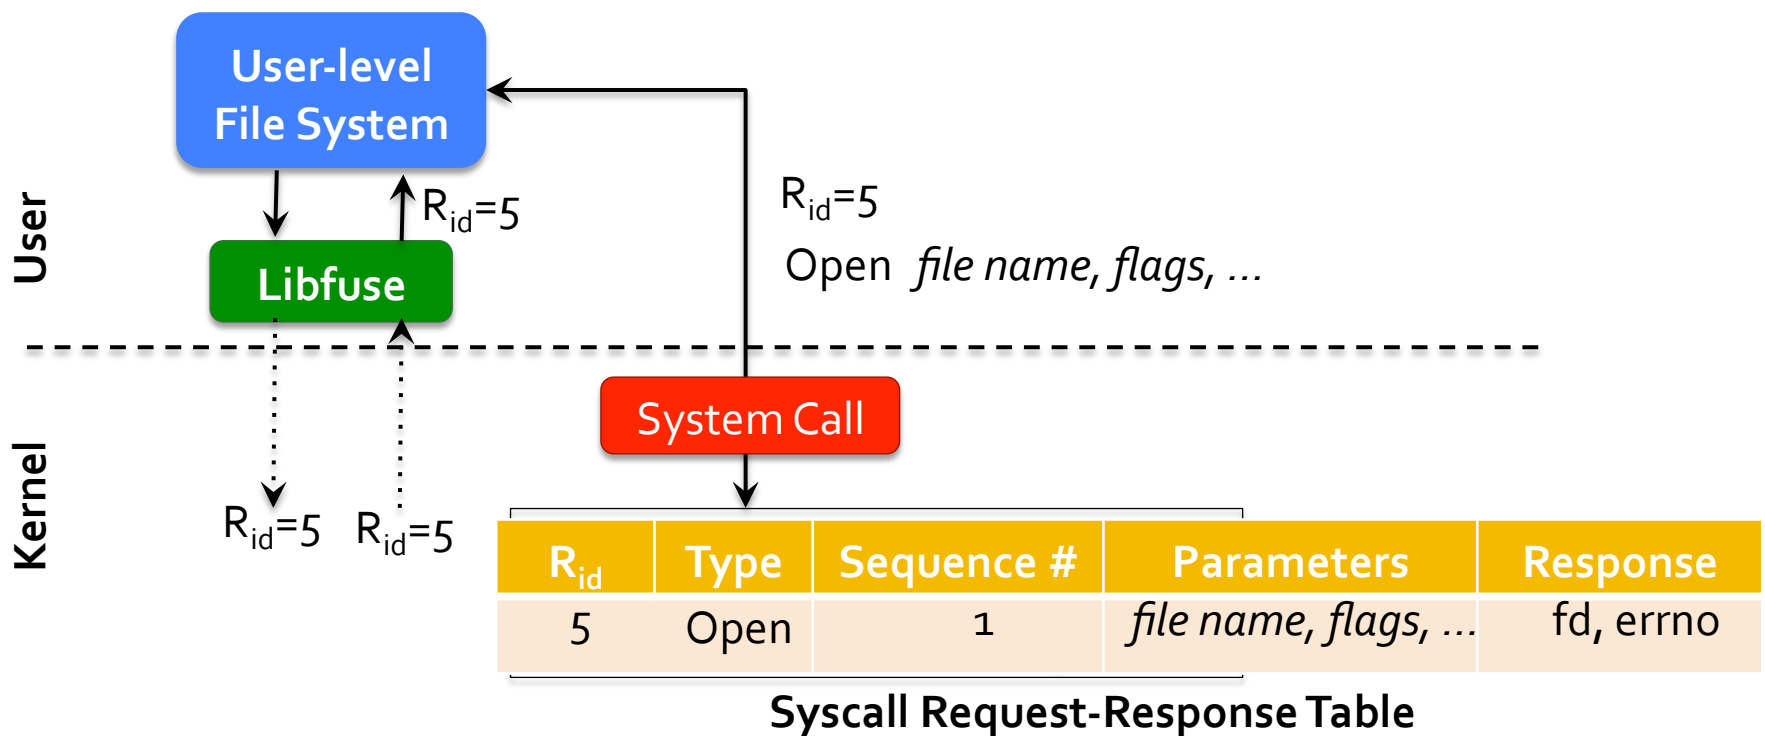
Request Tagging

- Correlates FS operation with external calls
 - Attach the fuse request id to the work thread




System Call Logging

- Track progress of individual operations
 - On replay: helps mimic execution of completed requests

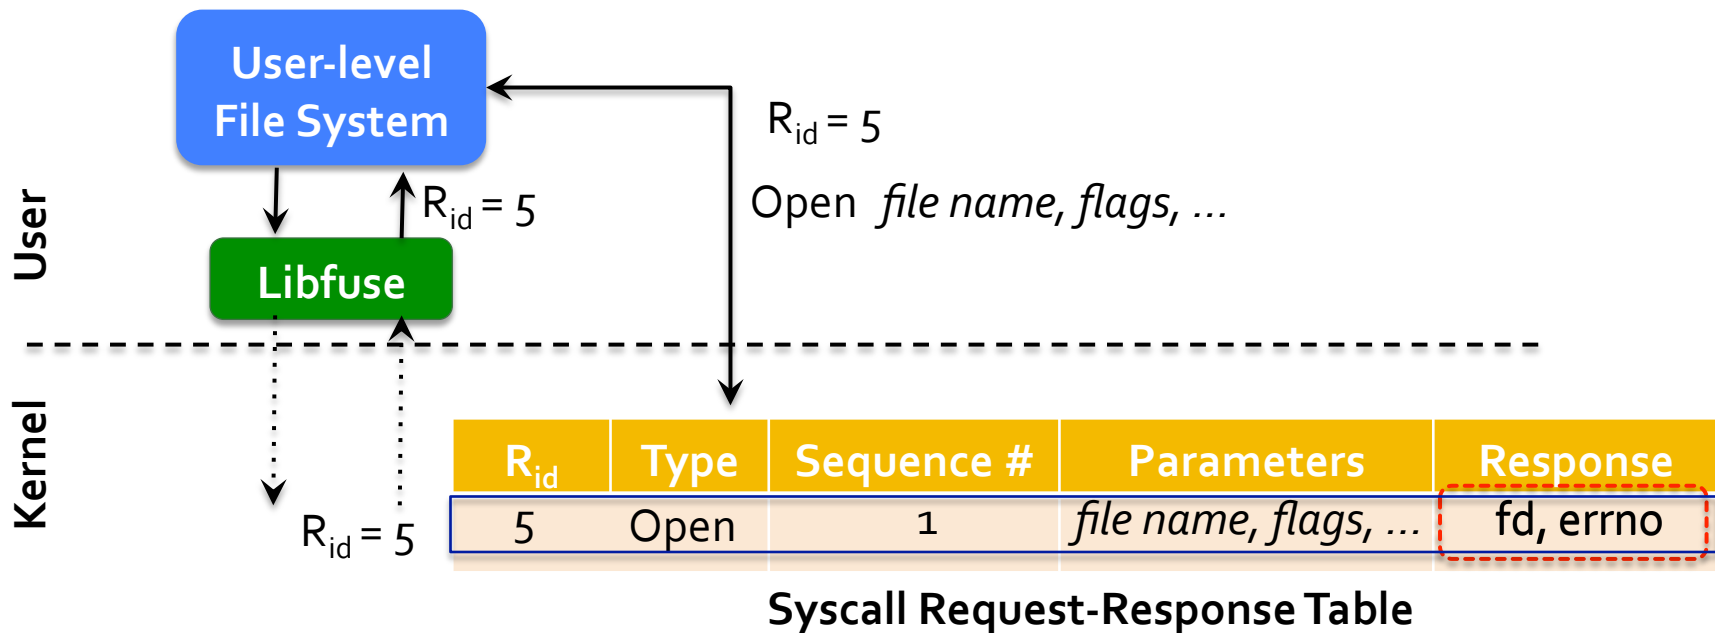


Components of Re-FUSE



Recovery

- Restore FS state needed to execute requests
 - Leverage cached state at the VFS layer
- Re-execute in-flight requests on restart
 - Leverage request queue at the KFM layer



Outline

- Background
- Challenges
- Re-FUSE
- Evaluation
- Conclusions

Evaluation

- Address the following questions
 - Implementation effort to work with Re-FUSE?
 - Can Re-FUSE hide failures from applications?
 - Performance overheads during user workloads?
- Experimental setup
 - 2.2 GHz Opteron, 2GB Ram, 2 80GB WDC disk
 - Implemented in Linux 2.6.18, FUSE (2.7.4)
 - Re-FUSE: 3300 loc in Linux kernel, 1000 loc in FUSE

Generality of Re-FUSE

	File System	Original	Added	Modified
Block-based interface	NTFS-3g	32K	10	1
Pass through interface	AVFS	39K	4	1
Network-based interface	SSHFS	4K	3	2

Code changes in individual FSes

- Additions
 - Daemonize user-level process
 - Notify local state (such as external file handles)
- Modifications
 - New mount interface that includes restart flag as parameter


Robustness of Re-FUSE

- Inject transient faults
 - Crashes in FS
- Inspect appl., FS state
 - ✗ - bad ✓ - good
- NTFS-3g results
 - SSHFS, AVFS in paper
- Random fault injection
 - 100 faults for all 3 FSES

		FUSE NTFS-3g			Re-FUSE NTFS-3g		
Operation	NTFS-3g Function	Application?	FS Consistent?	FS Usable?	Application?	FS Consistent?	FS Usable?
create	fuse_create	✗	✗	✗	✓	✓	✓
mkdir	fuse_create	✗	✗	✗	✓	✓	✓
symlink	fuse_create	✗	✗	✗	✓	✓	✓
link	link	✗	✗	✗	✓	✓	✓
rename	link	✗	✗	✗	✓	✓	✓
open	fuse_open	✗	✓	✗	✓	✓	✓
						✓	✓
						✓	✓

Re-FUSE successfully hides failures

NTFS-3g Performance Overheads




Outline

- Overview
- Challenges
- Re-FUSE
- Evaluation
- Conclusions

Conclusions

"Failure is not falling down but refusing to get up."

- Chinese proverb

- Reliability through restartability
 - A generic mechanism to restart user-level file systems
- **Principle:** inconsistent  consistent state
 - Inconsistency due to incomplete operations
 - Track progress of operations to continue from last execution
- Novel techniques
 - *Request tagging*: differentiate between serviced requests
 - *System call logging*: tracks sequence of operations
 - *Non-interruptible system call*: move threads to safe state
 - *Page versioning*: minimize logging overheads

Thank You!

Questions?



Advanced Systems Lab (ADSL)
University of Wisconsin-Madison
<http://www.cs.wisc.edu/adsl>

